

Computer Architecture and System Programming Laboratory

TA Session 5

Listing file
LOOP, EQU
Local labels
System calls

Producing a listing file: > nasm -f elf sample.s -l sample.lst



```
1
2          section .data
3 00000000 78563412      numeric:    DD 0x12345678
4 00000004 616263      string:      DB 'abc'
5 00000007 00000000      answer:      DD      0
6
7          section .text
8          global _start
9
10         _start:
11
12 00000000 60          pushad
13 00000001 6A02      push dword 2
14 00000003 6A01      push dword 1
15 00000005 E815000000 CALL myFunc
16
17 0000000A A307000000      mov [answer], eax
18 0000000F 83C408      add esp, 8
19 00000012 61          popad
20 00000013 BB00000000      mov ebx, 0
21 00000018 B801000000      mov eax, 1
22 0000001D CD80          int 0x80
23
24         myFunc:
25 0000001F 55          push    ebp
26 00000020 89E5      mov     ebp, esp
27 00000022 8B4508      mov    eax, dword [ebp+8]
28 00000025 8B5D0C      mov    ebx, dword [ebp+12]
29
30         myFunc_code:
31         returnFrom_myFunc:
32 0000002A 89EC      mov     esp, ebp
33 0000002C 5D          pop     dword ebp
34 0000002D C3          ret
```

first column is line number in the listing file

second column is relative address of where code will be placed in memory
- each section starts at relative address 0

third column is compiled code

forth column is original code

'CALL myFunc' is compiled to
opcode E8 followed by a **4-byte target address**, relative to the next instruction after the call.

→ address of myFunc label = 0x1F

→ address of the next instruction after the call (i.e. 'mov [answer], eax') is 0xA

→ $0x1F - 0xA = 0x15$, and we get exactly the binary code written here 'E815000000'

0x15 is how many bytes EIP should jump forward

```
Breakpoint 1, 0x08048080 in _start ()      executable
(gdb) x /1xg $eip
0x8048080 <_start>:      0x0015e8016a026a60
(gdb)
0x8048088 <_start+8>:      0x83080490b7a30000
(gdb)
0x8048090 <returnAddress+6>:      0x00000000bb6108c4
(gdb)
0x8048098 <returnAddress+14>:      0x5580cd00000001b8
(gdb)
0x80480a0 <myFunc+1>:      0x0c5d8b08458be589
(gdb)
0x80480a8 <myFunc_code>:      0x0000c35dec89d801
```

Advanced Instructions - LOOP

LOOP, LOOPE, LOOPZ, LOOPNE, LOOPNZ – loop with counter (CX or ECX)

```
mov ecx, 3
myLoop:
    ...
    loop myLoop, ecx    ; decrement ecx
                        ; if ecx != 0, jump to myLoop
```

LOOPE \equiv **LOOPZ**: jumps if the counter $\neq 0$ **and** ZF = 1

LOOPNE \equiv **LOOPNZ**: jumps if the counter $\neq 0$ **and** ZF = 0

Note: LOOP instruction does not set any flags

Note: if a counter is not specified explicitly, BITS setting dictates which is used

BITS directive specifies whether NASM should generate code designed to run on a processor operating in 16-bit mode, 32-bit mode, or 64-bit mode. The syntax is BITS 16, BITS 32, or BITS 64.

Local Labels Definition

A label beginning with a single period (.) is treated as a **local label**, which means that it is associated with the previous non-local label.

Example:

func1:

```
...  
.loop:  ←  
    dec eax  
    jne .loop
```

(this is indeed func1.loop)

func2:

```
...  
.loop:  ←  
    dec eax  
    jne .loop
```

(this is indeed func2.loop)

first character
can be: letter, _,
?, and .

valid characters in
labels are: letters,
numbers, _, \$, #, @,
~, ., and ?

Any local label can be
accessed from anywhere,
using the full, qualified
path through the relevant
non-local label

Each JNE instruction jumps to the closest .loop, because the two definitions of .loop are kept separate.

Linux System Calls

A system call is explicit request to the kernel, made via a software interrupt.

- System calls are low level functions, which represent the *interface* the kernel presents to user applications
- In Linux all low-level I/O is done by reading and writing files, regardless of what particular peripheral **device** is being accessed - a tape, a socket, even your terminal, they **are all files**.
- File is referenced by an integer file descriptor

Linux System Call format

- Put the system call number in EAX register
- Set up the arguments to the system call
 - The first argument goes in EBX, the second in ECX, then EDX, ESI, EDI, and finally EBP. If more than 6 arguments needed (not likely), the EBX register must contain the memory location where the list of arguments is stored.
- Call the relevant interrupt (for Linux it is 0x80)
- The result is usually returned in EAX

sys_read – read from a file

- system call number (in EAX): **3**
- arguments:
 - EBX: file descriptor (to read from it)
 - ECX: pointer to input buffer (to keep a read data into it)
 - EDX: maximal number of bytes to receive (maximal buffer size)
- return value (in EAX):
 - number of bytes received
 - On errors: negative number

```
section .bss
    buffer: resb 1

section .text
    global _start
_start:

    mov eax, 3           ; system call number (sys_read)
    mov ebx, 0           ; file descriptor (stdin)
    mov ecx, buffer      ; buffer to keep the read data
    mov edx, 1           ; bytes to read
    int 0x80             ; call kernel


    mov eax, 1           ; system call number (sys_exit)
    mov ebx, 0           ; exit status
    int 0x80             ; call kernel
```

sys_write – write into a file

- system call number (in EAX): 4
- arguments:
 - EBX: file descriptor (to write to it)
 - ECX: pointer to the first byte to write (beginning of the string)
 - EDX: number of bytes (characters) to write
- return value (in EAX):
 - number of bytes written
 - On errors: negative number

*\$ is a special label which exists **only at compile time**. \$ always points to the next free byte in the current section.*

By using \$ we can easily calculate the length of the message.



```
section .data
msg: db 'Hello',0xA      ; string to print
len: equ $ - msg         ; length of string

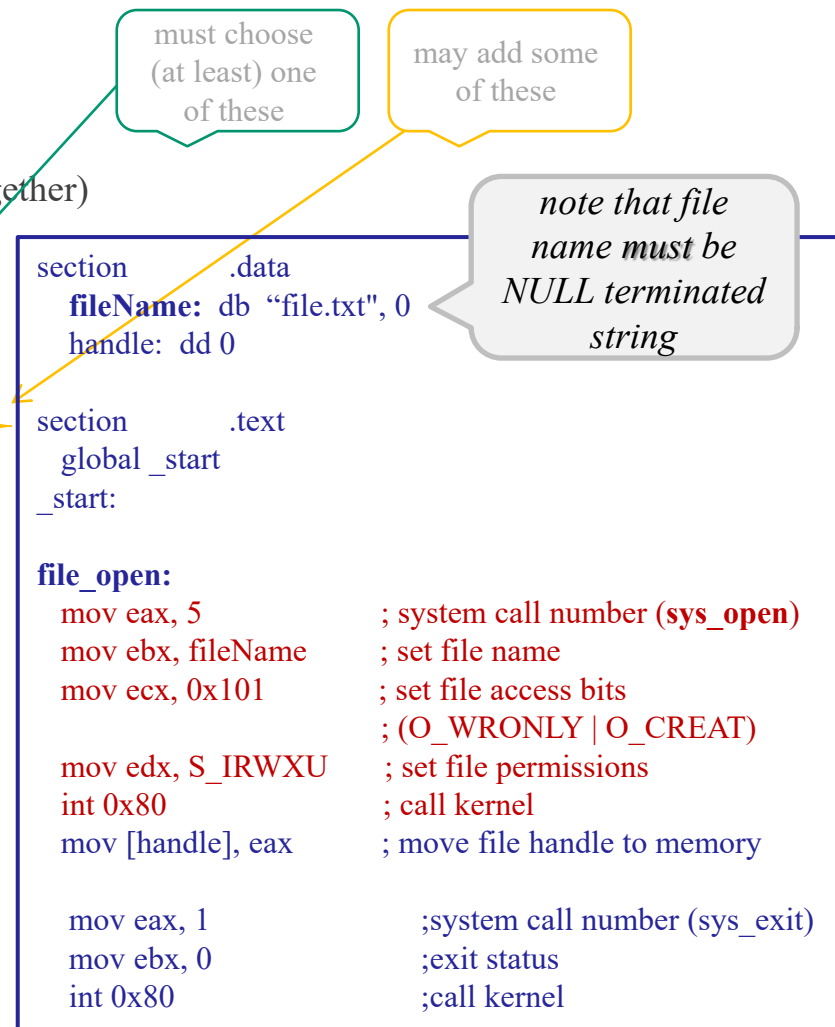
section .text
global _start
_start:

    mov eax,4             ;system call number (sys_write)
    mov ebx,1             ;file descriptor (stdout)
    mov ecx,msg           ;message to write
    mov edx,len           ;message length
    int 0x80              ;call kernel

    mov eax,1             ;system call number (sys_exit)
    mov ebx,0             ;exit status
    int 0x80              ;call kernel
```


sys_open - open a file

- system call number (in EAX): **5**
- arguments:
 - EBX: pathname of the file to open/create
 - ECX: set file access bits (can be bitwise OR'ed together)
 - O_RDONLY (0x000) open for reading only
 - O_WRONLY (0x001) open for writing only
 - O_RDWR (0x002) open for both reading and writing
 - O_APPEND (0x008) open for appending to the end of file
 - **O_CREAT (0x100)** create the file if it doesn't exist
 - EDX: set file permissions (in a case **O_CREAT** is set; can be bitwise OR'ed together)
 - S_IRWXU 0x700 ; **RWX** mask for owner
 - S_IRUSR 0x400 ; **R**(read) for owner **USR**(user)
 - S_IWUSR 0x200 ; **W**(write) for owner
 - S_IXUSR 0x100 ; **X**(execute) for owner
- return value (in EAX):
 - file descriptor
 - On errors: negative number



sys_lseek – change a file pointer

- system call number (in EAX): **19**
- arguments:
 - EBX: file descriptor
 - ECX: offset (number of bytes to move)
 - EDX: where to move from
 - SEEK_SET (0) ; beginning of the file
 - SEEK_CUR (1) ; current position of the file pointer
 - SEEK_END (2) ; end of file
- return value (in EAX):
 - Current position of the file pointer
 - On errors: SEEK_SET

```
section      .data
    fileName: db "file.txt", 0
    handle: dd 0

section      .text
    global _start
    _start:

file_open:
    mov eax, 5          ; system call number (sys_open)
    mov ecx, 0          ; set file access bits (O_RDONLY)
    mov ebx, fileName   ; set file name
    int 0x80            ; call kernel
    mov [handle], eax    ; move file handle to memory

    mov eax, 19         ; system call number (lseek)
    mov ebx, [handle]   ; file descriptor
    mov ecx, 15         ; number of byte to move
    mov edx, 0          ; move from beginning of the file
    int 0x80            ; call kernel

    mov eax, 1          ; system call number (sys_exit)
    mov ebx, 0          ; exit status
    int 0x80            ; call kernel
```

sys_close – close file

- system call number (in EAX): **6**
- arguments:
 - EBX: file descriptor
- return value (in EAX):
 - nothing meaningful
 - On errors: negative number

```
section      .data
    fileName: db "file.txt", 0
    handle : dd 0

section      .text
    global _start
    _start:

    file_open:
        mov eax, 5           ; system call number (sys_open)
        mov ecx, 0           ; set file access bits (O_RDONLY)
        mov ebx, fileName    ; set file name
        int 0x80             ; call kernel
        mov [handle], eax    ; move file handle to memory

        mov eax, 6
        mov ebx, [handle]
        int 0x80

        mov eax, 1           ;system call number (sys_exit)
        mov ebx, 0           ;exit status
        int 0x80             ;call kernel
```

Linux System Calls - Example

self reading

section .data

fileName: db "file.txt", 0

handle: dd 0

section .bss

buffer: resb 1

section .text

global _start

_exit:

mov ebx, [handle]

mov eax, 6 ; system call (sys_close)

int 0x80 ; call kernel

mov eax, 1 ; system call (sys_exit)

mov ebx, 0 ; exit status

int 0x80 ; call kernel

_start:

mov eax, 5 ; system call (sys_open)

mov ebx, fileName ; set file name

mov ecx, O_RDONLY ; set file access bits (O_RDONLY)

int 0x80 ; call kernel

mov [handle], eax ; move file handle to memory

_read:

mov eax, 3 ; system call (sys_read)

mov ebx, [handle] ; file handle

mov ecx, buffer ; buffer

mov edx, 1 ; read byte count

int 0x80 ; call kernel

cmp eax, 0

je _exit

mov eax, 4 ; system call (sys_write)

mov ebx, 1 ; stdout

mov ecx, buffer ; buffer

mov edx, 1 ; write byte count

int 0x80 ; call kernel

jmp _read