

# Introduction

*Amandyk Kartbayev*  
KBTU- 2021/2022

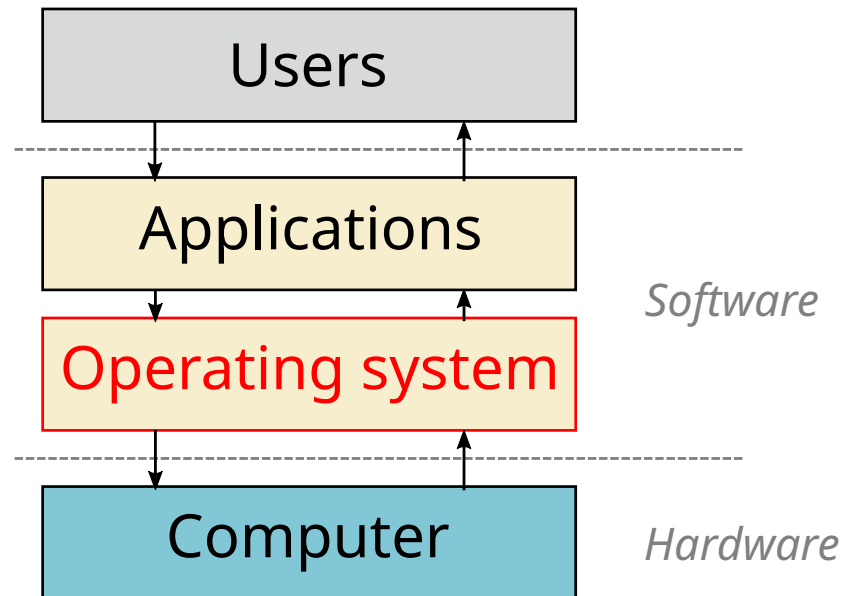
# Operating system definition

## Used to be a bit blurry

- Everything that was shipped by the operating system vendor
- But prone to abuse (e.g., US vs Microsoft, 98)

## (Somewhat) clearer definition

An operating system (OS) is the layer of software that manages a computer's resources for its users and their applications.




# Computer organization

## Types of computers



Desktop computer



Blade server (by Dmitry Nosachev - CC BY-SA 4.0)



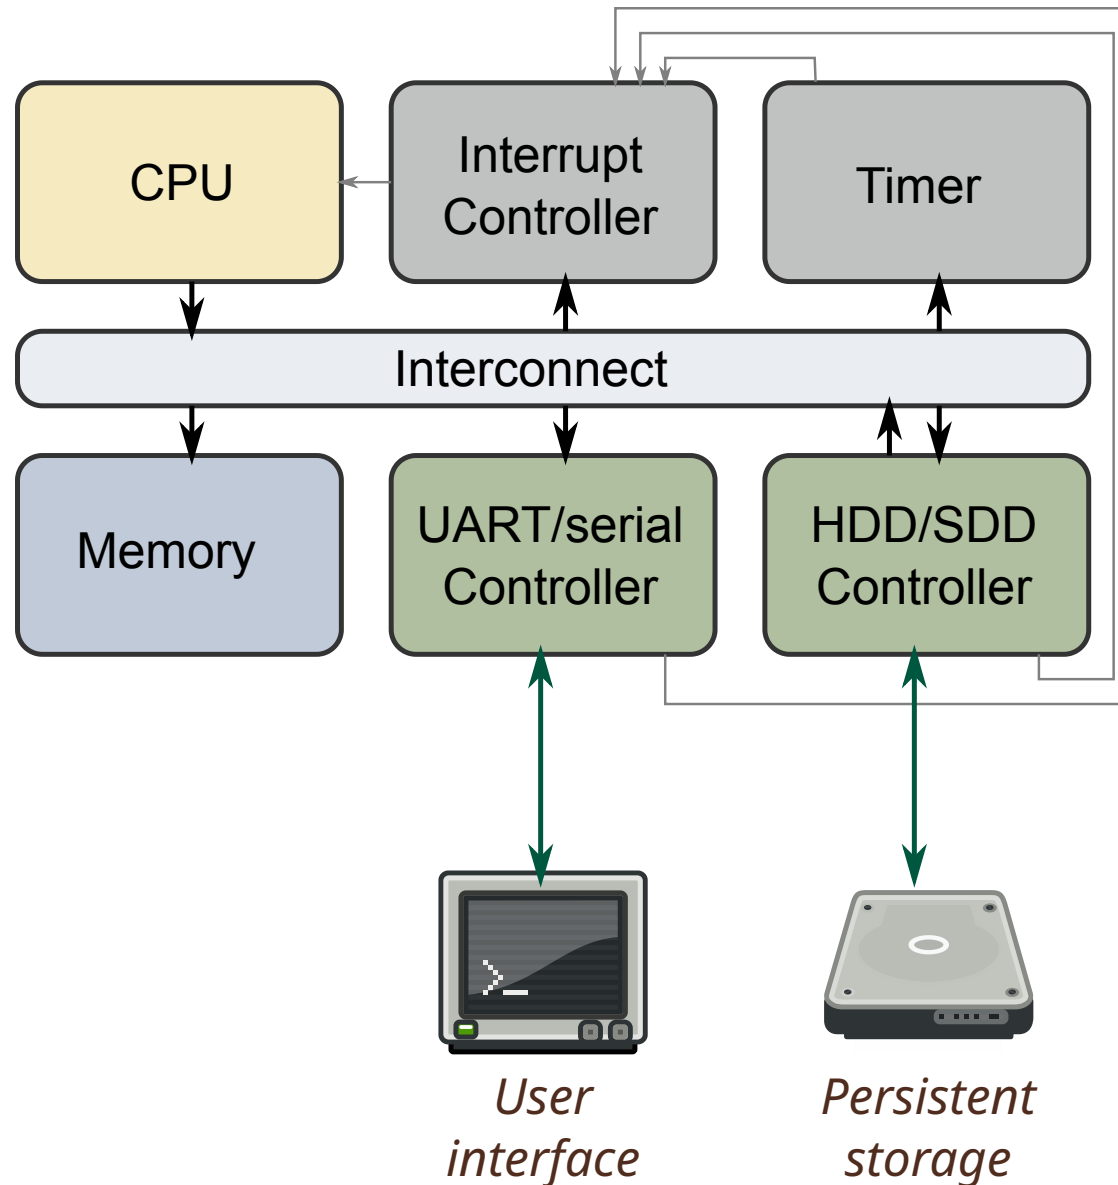
Internet of Things (by Miiichiaiel Hieinizilieir, CC BY-SA 4.0)



Car computer

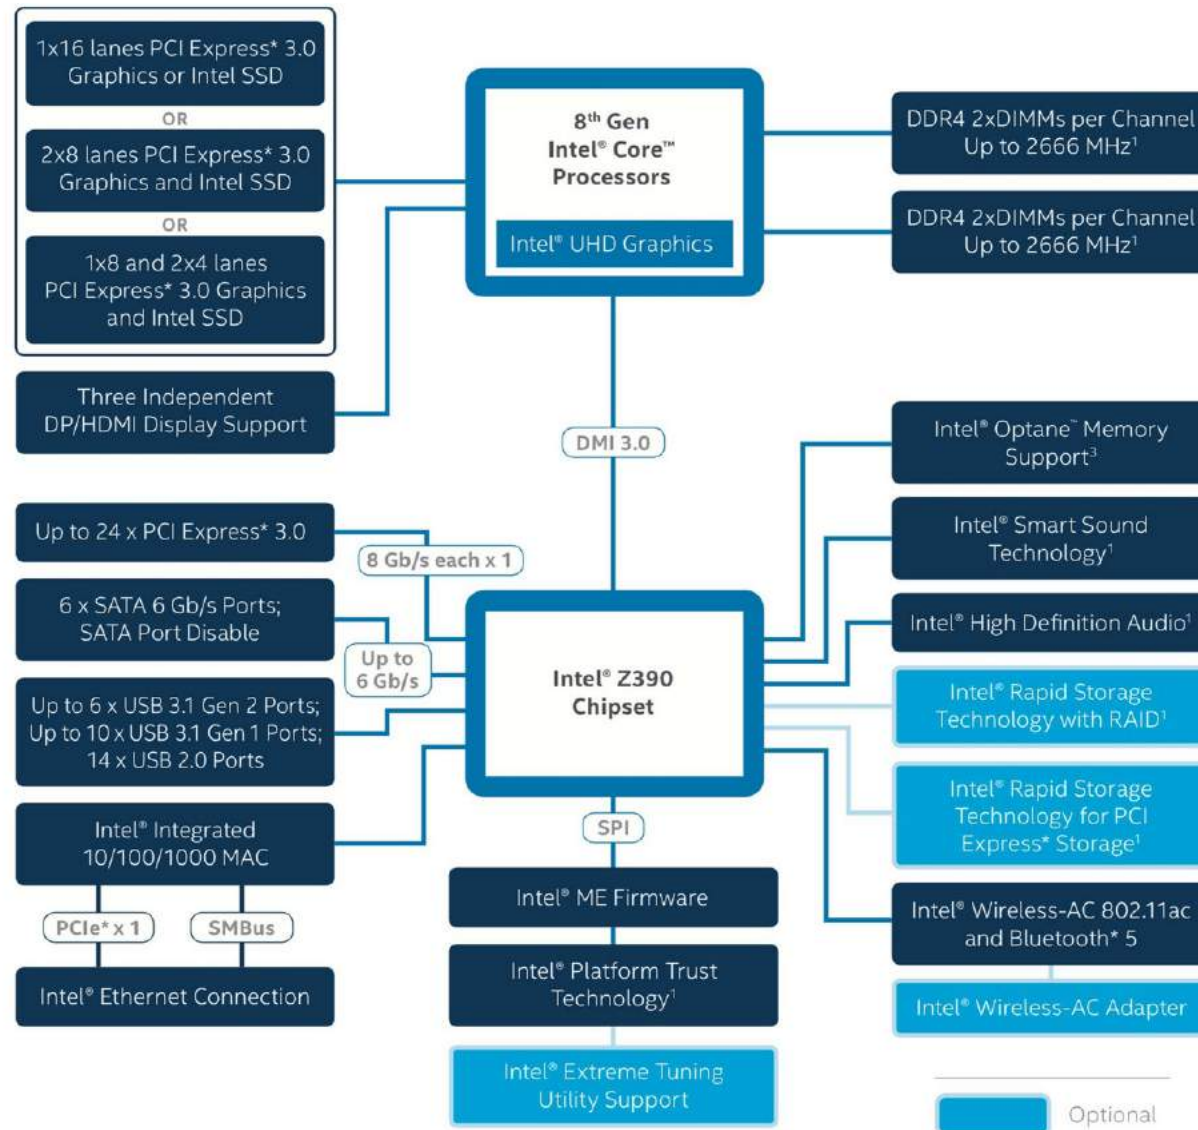
# Computer organization

## The bare minimum



# Computer organization

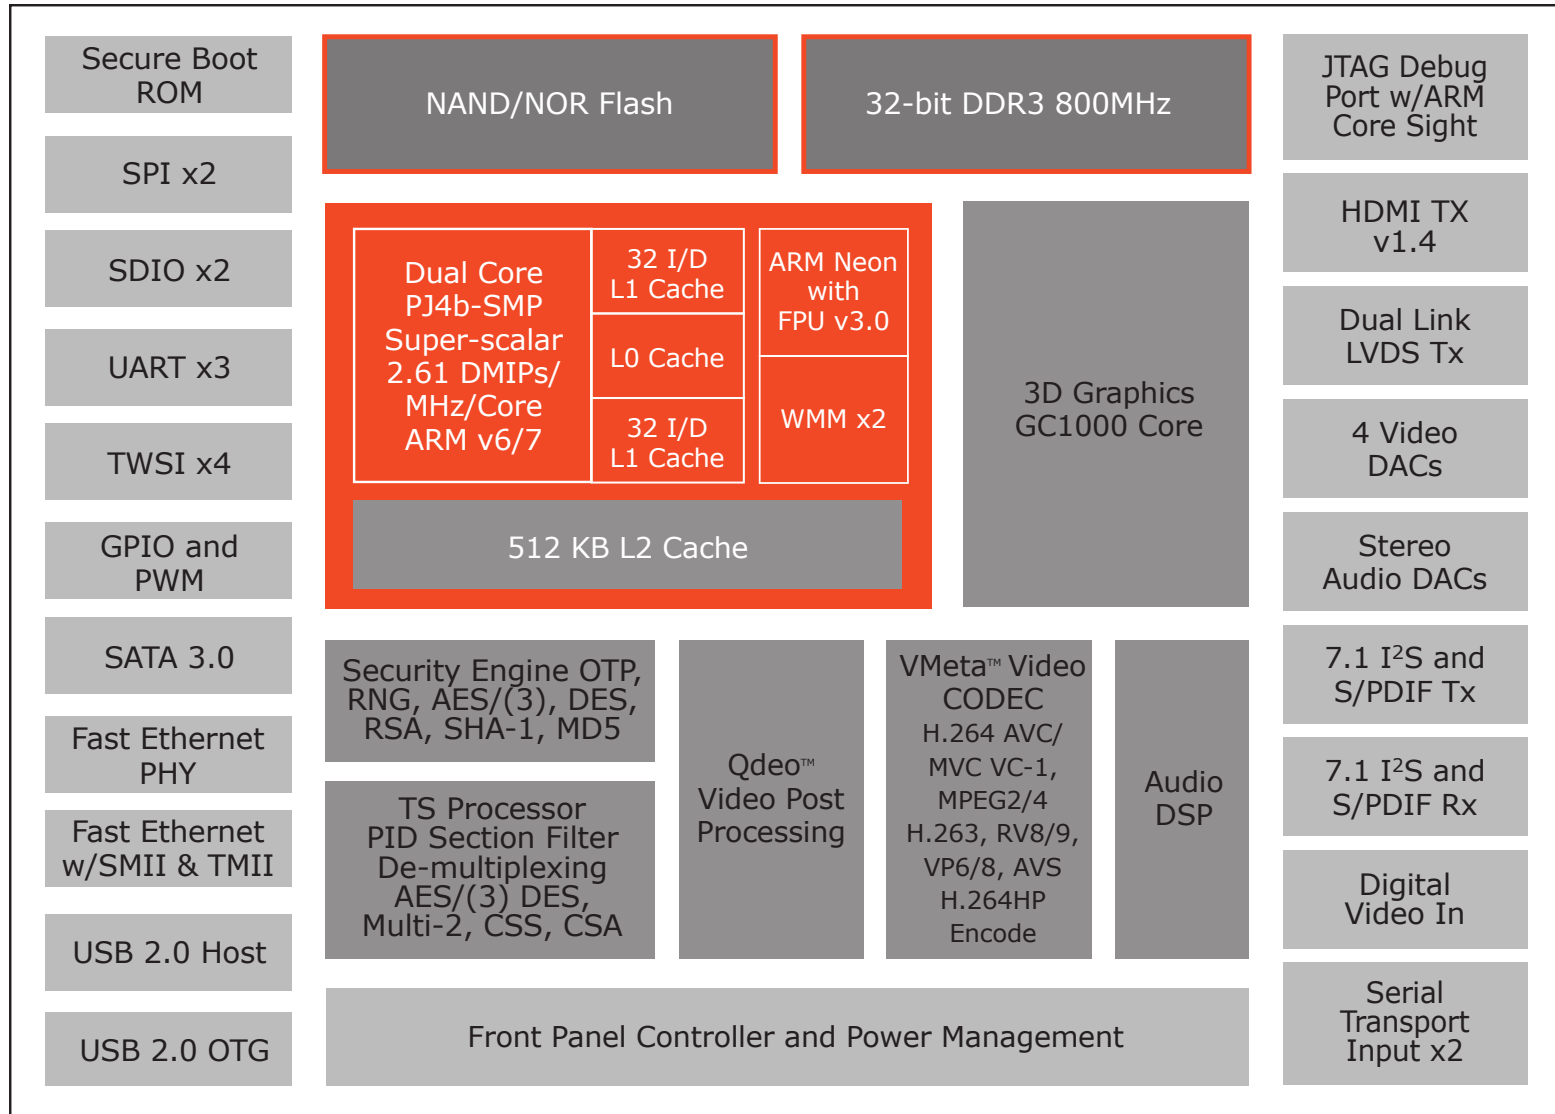
## General-purpose computing



Intel Chipset Z390 (2018)

# Computer organization

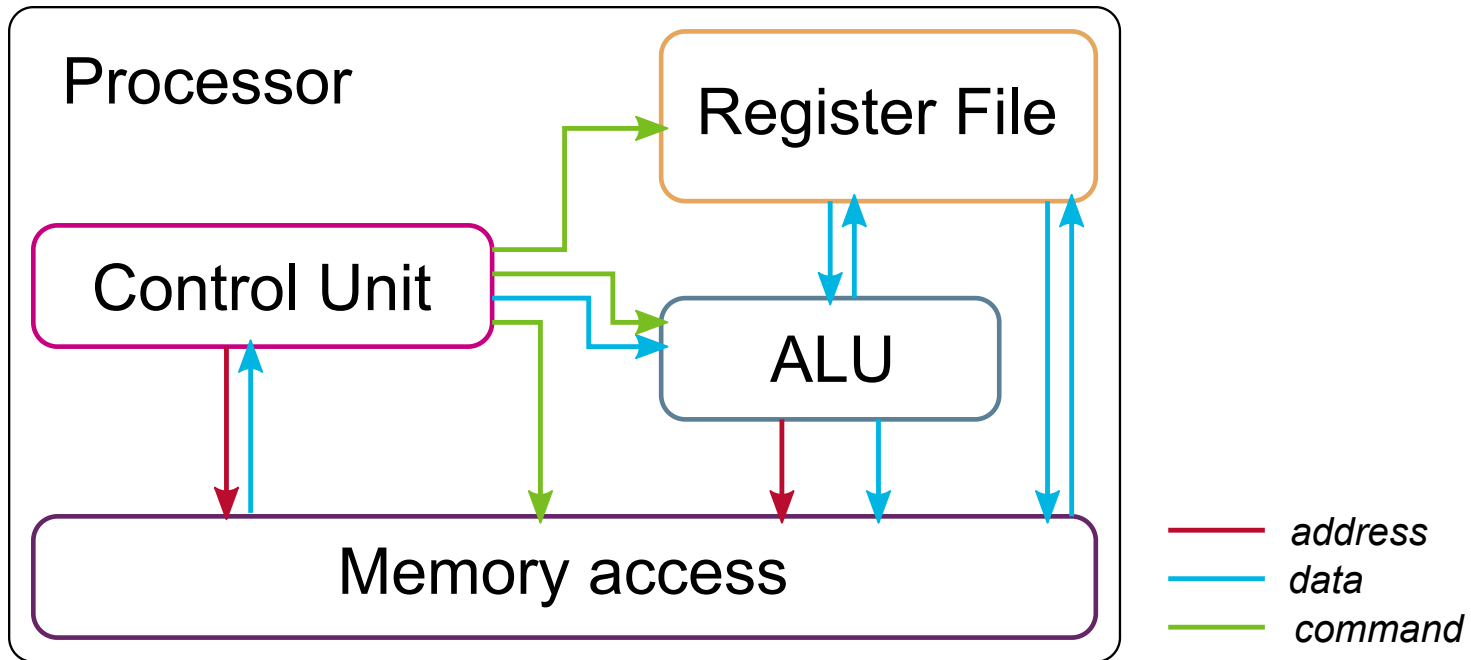
## Embedded computing



Marvell Armada 1500 (similar to SoC found in Google Chromecast)

# Hardware components

## Processor




- Fetches instructions from memory, decodes them, and executes them
- Characterized by an *Instruction Set* (e.g., i686, x86\_64, AArch64)
  - Access to memory, arithmetic/logical operations, control flow
- Contains a set of *registers*
  - General-purpose registers, program counter, status register

# Hardware components

## Memory

- Set of **addressable** bytes that can hold numerical values



- Organized in a hierarchy of layers (with various access times)

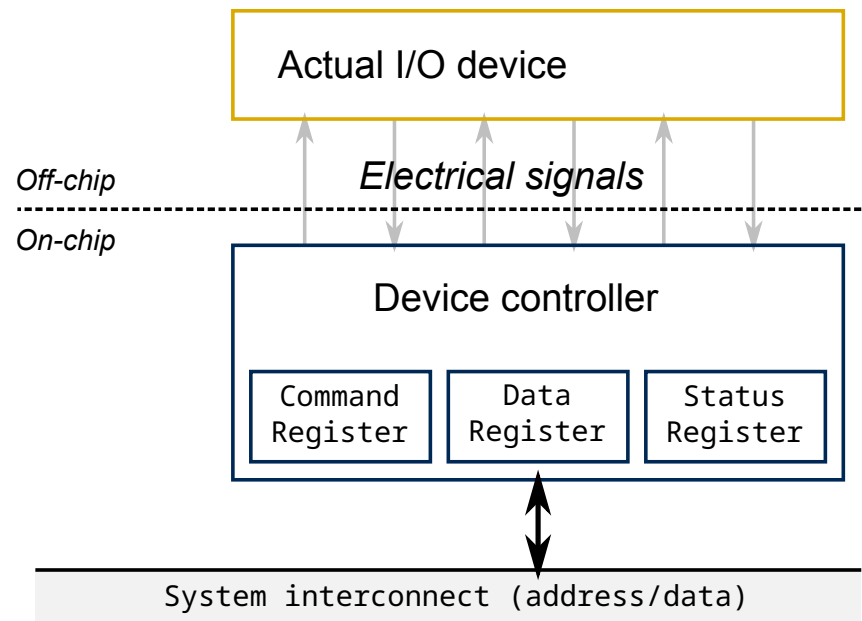
Memory layer	Access time	Accessibility
CPU registers	Immediate	Via CPU instructions
CPU cache	Few cycles	Not addressable (transparent)
RAM/Main memory	Few hundred cycles	Addressable from CPU
Second-level storage	Few thousand cycles	Indirectly addressable (see later in course)



# Hardware components

## I/O Devices

- Controllers connected to system interconnect
- Devices connected to controllers
- The controller provides an interface for accessing the device resources/functionalities
  - Via device registers



- Memory-mapped access
  - Device registers mapped into memory address space
  - Accessible through regular memory instructions
- Port-mapped access
  - Device registers mapped into special I/O space
  - Accessible through special I/O instructions

# Hardware components

---

## Interconnects (buses, networks)

- Transfer data between components
- Characterized by various features, speed, bandwidth, etc.

### CPU to memory buses

- Cache bus (aka Back side bus in Intel lingo)
- Memory bus (aka Front side bus in Intel lingo)
  - Now implemented by AMD HyperTransport, or Intel QuickPath Interconnect

### CPU to device buses

Bus	Created	Bandwidth	Type
ISA	1981	~8 MiB/s	Expansion
IDE	1986	~8 MiB/s	Mass-storage
PCI	1992	~133 MiB/s	Expansion
AGP	1997	~266 MiB/s	Video card
SATA 1.0	2000	~150 MiB/s	Mass-storage
PCI-e 1.x	2003	~250 MiB/s per lane	Expansion/Video card
<b>PCI-e 5.x</b>	<b>2019</b>	<b>~8 GiB/s per lane</b>	<b>Expansion/Video card</b>

# OS definition, part II

---

## Roles

In order to manage a computer's resources, an OS needs to play various roles:

- Referee
- Illusionist
- Glue

## Design principles

A well-constructed OS needs to achieve various design goals:

- Reliability
- Security
- Portability
- Performance

# OS roles

---

## Referee

Manage the **resource sharing** between applications

- Resource allocation (e.g., CPU, memory, I/O devices)
- Isolation (e.g., fault isolation)
- Communication (e.g., safe communication)



## Illusionist

Abstraction of hardware via **resource virtualization**

- Mask scarcity of physical resources
- Mask potential hardware failure



## Glue

Set of **common services** to applications

- Hardware abstraction
- Filesystem, message passing, memory sharing



# OS design principles

---

## Reliability

- OS does exactly what it is designed to do
- Related to *availability*: OS is always usable

## Security

- OS is secure if cannot be compromised by a malicious attacker
- Related to *privacy*: data is only accessible by authorized users
- Enforcement mechanisms vs security policies

## Portability

- OS provides the same abstractions regardless of the underlying hardware
  - Applications
  - System libraries
  - Kernel

## Performance

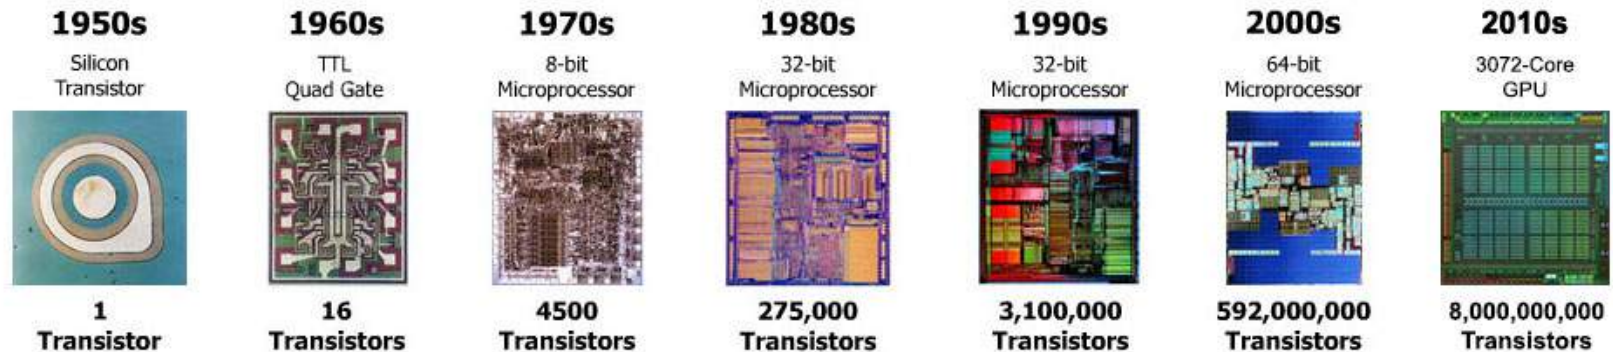
- Overhead
- Fairness
- Latency
- Throughput
- Predictability

*Binary compatibility is **so** important that I do not want to have anything to do with kernel developers who don't understand that importance. [...] The **only** reason for an OS kernel existing in the first place is to serve user-space.*  
*Linus Torvalds, LKML, 2012*

# OS history

## Three major phases

- Hardware is very expensive (1955-65)
- Hardware slowly becomes affordable (1965-80)
- Hardware becomes dirt cheap (1980-present)



# OS history

## Hardware is very expensive (1955-65)

- Introduction of the transistor in the mid-50s
- Expensive mainframes, operated by humans
- Read from punch card, run job, print result
- Batch systems in order to better serialize jobs



IBM 7090, circa 1960

## Software

- OS is simply a runtime *library* (common I/O functions)
- Applications have full access to hardware

# OS history

## Hardware slowly becomes affordable (1965-80)

- Use of integrated circuits
- Mainframe computers
  - E.g., IBM System/360
- Personal computers
  - E.g., DEC PDP-11



## Software

- IBM OS/360
  - Multiprogramming, memory protection
- Multics
  - Timesharing, dynamic linking, security, hierarchical file-system
- UNIX, BSD/SystemV variants
  - C language
- POSIX



# OS history

---

## Hardware becomes dirt cheap (1980-present)

- Continuation of Moore's law
- Personal computers
  - Command line interfaces
  - GUI
- Pervasive computers
  - Desktops, laptops
  - Smartphones, tablets
  - Embedded systems
  - Data centers



# OS history

## Progression over time

Metric	1981	1997	2014	Factor(2014/1981)
CPU performance (MIPS)	1	200	2500	2.5 K
CPU/computer	1	1	10+	10+
MIPS cost	\$100K	\$25	\$0.20	500 K
DRAM (MiB)/\$	\$500	\$0.5	\$0.001	500 K
Disk (GiB)/\$	\$333	\$0.14	\$0.00004	8 M
WAN (bps)	300	256K	20M	60 K
LAN (bps)	10M	100M	10G	1000+
Ratio users to computers	100:1	1:1	1:several	100+

Numbers taken from OSPP textbook

# Future of OSes

## End of Moore's law

- Make better use of the same area
  - Deep redesign of processors?
- 3-D stacking
  - Multiprocessors
  - Multi- and many-cores
- Quantum computing



## From the very small... to the super big

- Power-efficient IoT devices
  - Smart home, smart city, smart [blank]
- Giant data centers
  - Very large-scale storage

## Heterogeneity

- Different processors on same chip
  - E.g., ARM big.LITTLE, Intel Hybrid Technology
- Specialized computing accelerators
  - GP-GPUs, FPGAs, AI accelerators

# Manual approach

## Code example

### main.c

```
#include <stdio.h>
#include <stdlib.h>

#include "fact.h"

int main(int argc, char **argv)
{
    int n;
    if (argc < 2) {
        fprintf(stderr,
            "Usage: myfact number\n");
        exit(1);
    }
    n = atoi(argv[1]);
    printf("fact(%d) = %d\n",
        n, fact(n));
    return 0;
}
```

main.c

### fact.h

```
#ifndef FACT_H_
#define FACT_H_

int fact(int n);

#endif /* FACT_H_ */
```

fact.h

### fact.c

```
#include "fact.h"

int fact(int n) {
    if (n == 0)
        return 1;
    return n * fact(n - 1);
}
```

fact.c

### README.md

#### # Overview

This program computes the factorial of a number

README.md

# Manual approach

---

## Compilation

```
$ gcc -Wall -Wextra -Werror -c -o fact.o fact.c
$ gcc -Wall -Wextra -Werror -c -o main.o main.c
$ gcc -Wall -Wextra -Werror -o myfact main.o fact.o
```

```
$ ./myfact
Usage: myfact number
$ ./myfact 5
fact(5) = 120
```

```
$ pandoc -o README.html README.md
$ firefox README.html
```

## On the long run...

Now, what if:

- fact.c changes? main.c changes? fact.h changes?
- I want to change the compilation options?
- I want to recompile this code on another computer?
- I want to share this code?

Solution is to **automate the build process!**

# Introduction

---

## Definition

A *Makefile* is a file containing a set of rules used with the *make* build automation tool.

The two following commands are equivalent:

```
$ ls  
Makefile ...  
$ make  
$ make -f Makefile
```

The set of Makefile rules usually represents the various steps to follow in order to build a program: it's the building *recipe*.

# Introduction

---

## Anatomy of a rule

**target:** [list of prerequisites]

[ <tab> command ]

- For `target` to be generated, the prerequisites must all exist (or be generated if necessary)
- `target` is generated by executing the specified `command`
- `target` is generated only if it does not exist, or if one of the prerequisites is more recent
  - Prevents from building everything each time, but only what is necessary

## Commenting

- Lines prefixed with `#` are not evaluated

```
# This is a comment
```

# Version 0.1

---

## Basic rules

**myfact: main.o fact.o**

```
gcc -Wall -Wextra -Werror -o myfact main.o fact.o
```

**main.o: main.c fact.h**

```
gcc -Wall -Wextra -Werror -c -o main.o main.c
```

**fact.o: fact.c fact.h**

```
gcc -Wall -Wextra -Werror -c -o fact.o fact.c
```

**README.html: README.md**

```
pandoc -o README.html README.md
```

Makefile\_v0.1

**\$ make**

```
gcc -c -o main.o main.c
```

```
gcc -c -o fact.o fact.c
```

```
gcc -o myfact main.o fact.o
```

**\$ make README.html**

```
pandoc -o README.html README.md
```



# Version 0.1

## all rule

```
all: myfact README.html
```

```
myfact: main.o fact.o
```

```
gcc -Wall -Wextra -Werror -o myfact main.o fact.o
```

```
main.o: main.c fact.h
```

```
gcc -Wall -Wextra -Werror -c -o main.o main.c
```

```
fact.o: fact.c fact.h
```

```
gcc -Wall -Wextra -Werror -c -o fact.o fact.c
```

```
README.html: README.md
```

```
pandoc -o README.html README.md
```

Makefile\_v0.1

```
$ make
```

```
gcc -c -o main.o main.c
```

```
gcc -c -o fact.o fact.c
```

```
gcc -o myfact main.o fact.o
```

```
pandoc -o README.html README.md
```

# Version 0.1

---

## clean rule

```
all: myfact README.html
...
clean:
    rm -f myfact README.html main.o fact.o
```

Makefile\_v0.1

```
$ make
gcc -c -o main.o main.c
gcc -c -o fact.o fact.c
gcc -o myfact main.o fact.o
pandoc -o README.html README.md
```

```
$ make clean
rm -f myfact README.html main.o fact.o
```

# Version 0.1

---

## A first and basic Makefile

```
all: myfact README.html

myfact: main.o fact.o
    gcc -Wall -Wextra -Werror -o myfact main.o fact.o

main.o: main.c fact.h
    gcc -Wall -Wextra -Werror -c -o main.o main.c

fact.o: fact.c fact.h
    gcc -Wall -Wextra -Werror -c -o fact.o fact.c

README.html: README.md
    pandoc -o README.html README.md

clean:
    rm -f myfact README.html main.o fact.o
```

Makefile\_v0.1

- Was good enough for Project #1
  - *(No need to generate html out of markdown --pandoc is not installed on CSIF, and also it's just for the example)*

# Version 1.0

---

## How to avoid redundancy...?

*A good programmer is a lazy programmer!*

```
all: myfact README.html
```

```
myfact: main.o fact.o
```

```
gcc -Wall -Wextra -Werror -o myfact main.o fact.o
```

```
main.o: main.c fact.h
```

```
gcc -Wall -Wextra -Werror -c -o main.o main.c
```

```
fact.o: fact.c fact.h
```

```
gcc -Wall -Wextra -Werror -c -o fact.o fact.c
```

```
README.html: README.md
```

```
pandoc -o README.html README.md
```

```
clean:
```

```
rm -f myfact README.html main.o fact.o
```

Makefile\_v0.1

# Version 1.0

---

## Automatic variables in commands

- `$@`: replaced by name of target
- `$<`: replaced by name of **first** prerequisite
- `$^`: replaced by names of **all** prerequisites

```
all: myfact README.html
```

```
myfact: main.o fact.o
```

```
gcc -Wall -Wextra -Werror -o $@ $^
```

```
main.o: main.c fact.h
```

```
gcc -Wall -Wextra -Werror -c -o $@ $<
```

```
fact.o: fact.c fact.h
```

```
gcc -Wall -Wextra -Werror -c -o $@ $<
```

```
README.html: README.md
```

```
pandoc -o $@ $<
```

```
clean:
```

```
rm -f myfact README.html main.o fact.o
```

# Version 1.0

---

## Pattern rules

A pattern rule `%.o: %.c` says how to generate *any* file `<file>.o` from another file `<file>.c`.

```
all: myfact README.html
```

```
myfact: main.o fact.o
```

```
gcc -Wall -Wextra -Werror -o $@ $^
```

```
%.o: %.c fact.h
```

```
gcc -Wall -Wextra -Werror -c -o $@ $<
```

```
%.html: %.md
```

```
pandoc -o $@ $<
```

```
clean:
```

```
rm -f myfact README.html main.o fact.o
```

# Version 1.0

## Variables

```
CC      := gcc
CFLAGS  := -Wall -Wextra -Werror
CFLAGS += -g
PANDOC  := pandoc
```

```
all: myfact README.html
```

```
myfact: main.o fact.o
```

```
$(CC) $(CFLAGS) -o $@ $^
```

```
%.o: %.c fact.h
```

```
$(CC) $(CFLAGS) -c -o $@ $<
```

```
%.html: %.md
```

```
$(PANDOC) -o $@ $<
```

```
clean:
```

```
rm -f myfact README.html \
    main.o fact.o
```

```
$ make
gcc -Wall -Wextra -Werror -g -c -o main.o main.c
gcc -Wall -Wextra -Werror -g -c -o fact.o fact.c
gcc -Wall -Wextra -Werror -g -o myfact main.o fact.o
pandoc -o README.html README.md
```

# Version 2.0

## More variables

```
targets := myfact README.html
objs    := main.o fact.o

CC      := gcc
CFLAGS  := -Wall -Wextra -Werror
CFLAGS += -g
PANDOC  := pandoc
```

```
all: $(targets)
```

```
myfact: $(objs)
```

```
$(CC) $(CFLAGS) -o $@ $^
```

```
%.o: %.c fact.h
```

```
$(CC) $(CFLAGS) -c -o $@ $<
```

```
%.html: %.md
```

```
$(PANDOC) -o $@ $<
```

```
clean:
```

```
rm -f $(targets) $(objs)
```

```
$ make
gcc -Wall -Wextra -Werror -g -c -o main.o main.c
gcc -Wall -Wextra -Werror -g -c -o fact.o fact.c
gcc -Wall -Wextra -Werror -g -o myfact main.o fact.o
pandoc -o README.html README.md
```



# Version 2.0

## Nice output

```
$ make
CC main.o
CC fact.o
CC myfact
MD README.html
```

```
$ make clean
CLEAN
```

...

```
myfact: $(objs)
    @echo "CC    @"
    @$(CC) $(CFLAGS) -o $@ $^

%.o: %.c fact.h
    @echo "CC    @"
    @$(CC) $(CFLAGS) -c -o $@ $<

%.html: %.md
    @echo "MD    @"
    @$(PANDOC) -o $@ $<

clean:
    @echo "CLEAN"
    @rm -f $(targets) $(objs)
```

- In case of debug, how can we still see the commands that are executed?

# Version 3.0

## Conditional variables

```
...
```

```
ifneq ($(V),1)
```

```
Q = @
```

```
endif
```

```
myfact: $(objs)
```

```
    @echo "CC    $@"
```

```
    $(Q)$(CC) $(CFLAGS) -o $@ $^
```

```
%.o: %.c fact.h
```

```
    @echo "CC    $@"
```

```
    $(Q)$(CC) $(CFLAGS) -c -o $@ $<
```

```
%.html: %.md
```

```
    @echo "MD    $@"
```

```
    $(Q)$(PANDOC) -o $@ $<
```

```
clean:
```

```
    @echo "CLEAN"
```

```
    $(Q)rm -f $(targets) $(objs)
```

```
$ make
```

```
CC main.o
```

```
CC fact.o
```

```
CC myfact
```

```
MD README.html
```

```
$ make V=1
```

```
CC main.o
```

```
gcc -Wall -Wextra -Werror -g -c -o main.o main.c
```

```
CC fact.o
```

```
gcc -Wall -Wextra -Werror -g -c -o fact.o fact.c
```

```
CC myfact
```

```
gcc -Wall -Wextra -Werror -g -o myfact main.o fact.o
```

```
MD README.html
```

```
pandoc -o README.html README.md
```

# Version 3.0

## Generic rules vs dependency tracking

### Non-generic rule

```
%.o: %.c fact.h  
    @echo "CC   $@"  
    $(Q)$(CC) $(CFLAGS) -c -o $@ $<
```

### Generic rule

```
%.o: %.c  
    @echo "CC $@"  
    $(Q)$(CC) $(CFLAGS) -c -o $@ $<  
Makefile_v3.0
```

- How can we preserve the generic rule but also have accurate dependency tracking?

```
$ make  
CC main.o  
CC fact.o  
CC myfact  
MD README.html
```

```
$ make  
make: Nothing to be done for 'all'.
```

```
$ touch fact.h  
$ make  
make: Nothing to be done for 'all'.
```

# Version 3.0

## Rule composition

```
%.o: %.c  
    @echo "CC $@"  
    $(Q)$(CC) $(CFLAGS) -c -o $@ $<  
Makefile_v3.0
```

```
main.o: main.c fact.h  
fact.o: fact.c fact.h
```

- How can we have these additional rules be generated automatically and included in the Makefile?

```
$ make  
CC main.o  
CC fact.o  
CC myfact  
MD README.html
```

```
$ make  
make: Nothing to be done for 'all'.
```

```
$ touch fact.h  
$ make  
CC main.o  
CC fact.o  
CC myfact
```

# Version 3.0

---

## Use GCC for dependency tracking

```
#include <stdio.h>
#include "fact.h"

int fact(int n) {
    if (n == 0)
        return 1;
    return n * fact(n - 1);
}
```

```
$ gcc -Wall -Wextra -Werror -MMD -c -o fact.o fact.c
```

```
$ cat fact.d
fact.o: fact.c fact.h
```

# Version 3.0

## Dependency tracking Makefile integration

```
targets := myfact README.html
objs    := main.o fact.o
...
CFLAGS  := -Wall -Wextra -Werror -MMD
...
all: $(targets)

# Dep tracking *must* be below the 'all' rule
deps := $(patsubst %.o,%.d,$(objs))
-include $(deps)
...
%.o: %.c
    @echo "CC $@"
    $(Q)$(CC) $(CFLAGS) -c -o $@ $<
...
clean:
    @echo "clean"
    $(Q)rm -f $(targets) $(objs) $(deps)
```

Makefile\_v3.0

- `$(deps)` will be computed from `$(obj)` into `main.d fact.d`
- Prefix `-` ignores inclusion errors

# Version 3.0

---

## First run

- Dependency files don't exist but make won't complain
- GCC generates them

```
$ ls *.d
```

```
$ make  
CC main.o  
CC fact.o  
CC myfact  
MD README.html  
$ ls *.d  
main.d fact.d
```

```
$ cat main.d  
main.o: main.c fact.h  
$ cat fact.d  
fact.o: fact.c fact.h
```

## Following runs

- Dependency files are included by the Makefile
- They are used to compose the generic rule for object generation

```
$ make  
make: Nothing to be done for 'all'  
$ touch fact.h  
$ make  
CC main.o  
CC fact.o  
CC myfact
```

# Final Makefile

```
targets := myfact README.html
objs    := main.o fact.o

CC      := gcc
CFLAGS  := -Wall -Wextra -Werror -MMD
CFLAGS += -g
PANDOC  := pandoc

ifneq ($(V),1)
Q = @
endif

all: $(targets)

# Dep tracking *must* be below the 'all' rule
deps := $(patsubst %.o,%.d,$(objs))
-include $(deps)

myfact: $(objs)
    @echo "CC $@"
    $(Q)$(CC) $(CFLAGS) -o $@ $^

%.o: %.c
    @echo "CC $@"
    $(Q)$(CC) $(CFLAGS) -c -o $@ $<

%.html: %.md
    @echo "MD $@"
    $(Q)$(PANDOC) -o $@ $<

clean:
    @echo "clean"
    $(Q)rm -f $(targets) $(objs) $(deps)
```

Makefile\_v3.0



# GDB

---

## GNU project

- Started by Richard Stallman in 1983
- Free software, mass collaboration project in response to proprietary UNIX
  - Copyleft license: GNU GPL
  - User programs: text editor (Emacs), compiler (GCC toolchain), debugger (GDB), and various utilities (ls, grep, awk, make, etc.)
  - Kernel: GNU Hurd

## GDB

- GNU DeBugger
- Supports many languages
  - Including C and C++
- Inspection of program during execution
  - Execution flow
  - Data
- Helps finding errors like *segmentation fault*
- Read the fully-detailed manual:  
<https://sourceware.org/gdb/current/onlinedocs/gdb/>

# GDB usage

---

## Compilation flags

- Canonical compilation command line:

```
$ gcc [cflags] -o <output> <input>
```

- Optimize for speed (-O2)

```
$ gcc -Wall -Werror -O2 -o myprogram main.c
```

- Enable debugging support (-g)

```
$ gcc -Wall -Werror -g myprogram main.c
```

- Not recommended to use debugging along with optimizations
  - No optimization option is equivalent to -O0

# GDB usage

---

## Makefile digression

- During development, very useful to be able to debug your program
- For production, probably better to disable the debug support and activate all possible optimization support
  - Reduce size of the executable (can easily be by 50%!)
  - Increase performance (can also be by 50%!)

## Makefile automation

```
ifeq ($(D),1)
CFLAGS      += -g    # Enable debugging
else
CFLAGS      += -O2    # Enable optimization
endif
```

## Building mode

```
$ make D=1    # compile with debug support and no optimization
```

```
$ make        # compile with optimizations (production)
```

- Probably want to use `make clean` when changing building mode

# GDB usage

---

## Starting GDB

- Start GDB, specify the program to debug

```
$ gdb
...
(gdb) file myprogram
Reading symbols from myprogram...done.
(gdb)
```

- Or, start GDB with the program to debug as argument

```
$ gdb myprogram
...
Reading symbols from myprogram...done.
(gdb)
```

## Running the program

- Without any argument:

```
(gdb) run
```

- With arguments:

```
(gdb) run argv1 argv2...
```

# GDB usage

---

## Interactive help

- GDB offers an interactive shell
  - History management
  - Auto-complete (with TAB)

In order to discover what you can do, just ask:

```
(gdb) help
List of classes of commands:
aliases -- Aliases of other commands
breakpoints -- Making program stop at certain points
...

(gdb) help breakpoints
Making program stop at certain points.
List of commands:
awatch -- Set a watchpoint for an expression
break -- Set breakpoint at specified location
...

(gdb) help break
Set breakpoint at specified location.
break [PROBE_MODIFIER] [LOCATION] [thread THREADNUM] [if CONDITION]
...
```

# GDB usage

---

## Possible scenarios

### 1. Program doesn't have bugs:

- It will run fine until completion

```
$ ./myprogram  
I worked, hurray!
```

### 2. *Best-case* scenario, regarding bugs:

- *Segmentation fault*

```
$ ./myprogram  
segmentation fault (core dumped) ./myprogram
```

### 3. *Worst-case* scenario:

- Doesn't crash but wrong result

```
$ ./myprogram  
I work??, ??rray!
```

- Bugs that don't trigger any segmentation fault
- In this case, you'll probably have to spend more time...

# Segmentation faults

## Example #1

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

size_t foo_len (const char *s)
{
    return strlen(s);
}

int main (int argc, char *argv[])
{
    char *a = NULL;

    printf ("size of a = %d\n", foo_len(a));

    return 0;
}
```

## Execution

```
$ ./strlen-test
zsh: segmentation fault (core dumped) ./strlen-test
```

# Segmentation faults

## Run with GDB

- (After compiling the code with `-g`)

```
$ gdb ./strlen-test
(gdb) run
Starting program: /home/joel/tmp/test/strlen-test

Program received signal SIGSEGV, Segmentation fault.
0x00007ffff7abc446 in strlen () from /usr/lib/libc.so.6
(gdb)
```

## Backtrace

- First thing to do when getting a *segfault*:
  - Understand what is the sequence of calls that brought us there

```
(gdb) backtrace      # use just 'bt'
#0  0x00007ffff7abc446 in strlen () from /usr/lib/libc.so.6
#1  0x000000000040055e in foo_len (s=0x0) at strlen-test.c:7
#2  0x0000000000400583 in main (argc=1, argv=0x7fffffffd788) at strlen-test.c:14
```

## Investigate

- `foo_len()` is supposed to receive a pointer
- Here it receives `0` (aka `NULL`)
- Looks like this `NULL` pointer probably gets dereferenced in `strlen()`...



# Segmentation faults

---

Fix...

- Here, the problem is fairly obvious

```
size_t foo_len (const char *s)
{
    return strlen(s);
}

int main (int argc, char *argv[])
{
    char *a = "This is a valid string";

    printf ("size of a = %d\n", foo_len(a));

    return 0;
}
```

And, celebrate!

```
$ ./strlen-test
size of a = 22
```

# Segmentation faults

## Better fix

- Prevent the same bug from happening again

```
size_t foo_len (const char *s)
{
    assert(s && "String cannot be NULL here!");
    return strlen(s);
}

int main (int argc, char *argv[])
{
    char *a = NULL;

    printf ("size of a = %d\n", foo_len(a));

    return 0;
}
```

```
$ ./strlen-test
strlen-test: strlen-test.c:8: foo_len:
Assertion `s && "String cannot be NULL here!'" failed.
```

# Segmentation faults

## Example #2

```
#include<stdio.h>
#include<stdlib.h>

const static int len = 10;

int main(void)
{
    int *tab;
    unsigned int i;

    tab = malloc(len * sizeof(int));

    for (i = len - 1; i >= 0; i--)
        tab[i] = i;

    free(tab);
    return 0;
}
```

## Execution

```
$ ./tablen-test
segmentation fault (core dumped) ./tablen-test
```

# Segmentation faults

## Run GDB

```
$ gdb ./tablen-test
(gdb) run
Starting program: /home/joel/tmp/test/tablen-test

Program received signal SIGSEGV, Segmentation fault.
0x0000000000400535 in main () at tablen-test.c:14
14          tab[i] = i;
```

## Backtrace

```
(gdb) bt
#0  0x0000000000400535 in main () at tablen-test.c:14
```

- Except that here, it's not much of help...

## Inspect variables

- Display index `i` so that we know which index in the array was being accessed:

```
(gdb) print i
$1 = 4294967295
```

# Segmentation faults

---

Fix...

- Problem is a case of overflow
  - An unsigned int type automatically wraps from 0 to 4294967295

```
#include<stdio.h>
#include<stdlib.h>

const static int len = 10;

int main(void)
{
    int *tab;
    int i;

    tab = malloc(len * sizeof(int));

    for (i = len - 1; i >= 0; i--)
        tab[i] = i;

    free(tab);
    return 0;
}
```

# Tracking bugs

## Behavior bugs

- Behavioral bugs more complicated to find because program doesn't crash
- It's just that the output is wrong

```
#include <ctype.h>
#include <stdio.h>
#include <string.h>

int main(void)
{
    int i;
    char str[] = "Tracking bugs is my passion";

    printf("Before: %s\n", str);

    for (i = 0; i < strlen(str) - 1; i++)
        str[i] = toupper(str[i]);

    printf("After: %s\n", str);

    return 0;
}
```

## Execution

```
Before: Tracking bugs is my passion
After: TRACKING BUGS IS MY PASSION
```

# Tracking bugs

---

## Setting breakpoints

- Stop the program during the execution at a designated point
- Set as many breakpoints as necessary
- GDB will always stop the execution when reaching them

## Breaking at exact location in code

```
(gdb) break string-test.c:13  
Breakpoint 1 at 0x4005ef: file string-test.c, line 13.
```

```
(gdb) r  
Starting program: /home/joel/tmp/test/string-test  
Before: Tracking bugs is my passion
```

```
Breakpoint 1, main () at string-test.c:13  
13          str[i] = toupper(str[i]);
```

# Tracking bugs

## Breaking at a particular function

```
(gdb) b main
Breakpoint 1 at 0x40059f: file string-test.c, line 8.

(gdb) r
Starting program: /home/joel/tmp/test/string-test

Breakpoint 1, main () at string-test.c:8
8          char str[] = "Tracking bugs is my passion";
```

## Breaking only if condition is satisfied

```
(gdb) b string-test.c:13 if i == 5
Breakpoint 1 at 0x4005ef: file string-test.c, line 13.

(gdb) r
Starting program: /home/joel/tmp/test/string-test
Before: Tracking bugs is my passion

Breakpoint 1, main () at string-test.c:13
13          str[i] = toupper(str[i]);

(gdb) print i
$1 = 5
```



# Tracking bugs

## Dealing with breakpoints

- Set at least one breakpoint before running the program
  - Otherwise the program will run until completion
- Once the program stops and the gdb shell is available, a few options:
  1. Continue the execution until hitting the same or another breakpoint

```
(gdb) continue      # or just 'c'
```

2. Execute only the next line of code and break again

```
(gdb) step          # or just 's'
```

Careful, `step` enters function calls

3. Jump over function calls

```
(gdb) next          # or just 'n'
```

*Tip: typing `<enter>` in the interactive GDB shell repeats the last command*

# Tracking bugs

## Printing variables

```
int a = 2;  
char b = 'x';  
int *c = &a;  
char *s = "A string";  
  
... // <= breaking here
```

- Inspect the value of all your variables with command `print`

### Default

- By default, prints variables according to their type

```
(gdb) print a  
$1 = 2  
(gdb) p b  
$2 = 120 'x'  
(gdb) p c  
$3 = (int *) 0x7fffffffdd65c  
(gdb) p s  
$4 = 0x40070b "A string"
```

### Tweak

- Can tweak both the way `print` prints and what it prints

```
(gdb) print /x a  
$1 = 0x2  
(gdb) p /c b+2  
$2 = 122 'z'  
(gdb) p *c  
$3 = 2  
(gdb) p s[0]  
$4 = 65 'A'
```

# Tracking bugs

## Printing data structures

```
struct entry {  
    int    key;  
    char   *name;  
} obj = {  
    .key    = 2,  
    .name   = "toto",  
};  
  
struct entry *e = &obj;
```

- With `print`, you can access the pointer and the object it's pointing to:

```
(gdb) print e  
$1 = (struct entry *) 0x7fffffff640  
(gdb) print &obj  
$2 = (struct entry *) 0x7fffffff640  
(gdb) p *e  
$3 = {key = 2, name = 0x400734 "toto"}  
(gdb) p e->key  
$4 = 2  
(gdb) p obj.name  
$5 = 0x400734 "toto"
```

# Misc

---

## Setting watchpoint

- Breakpoints are for interrupting the execution flow at a specific location
- Watchpoints are for interrupting the program when a variable is modified

```
(gdb) watch i
Hardware watchpoint 2: i

(gdb) c
Continuing.
Before: Tracking bugs is my passion

Hardware watchpoint 2: i

Old value = 0
New value = 1
0x0000000000400612 in main () at string-test.c:12
12          for (i = 0; i < strlen(str) - 1; i++)
```

# Misc

## Other useful commands

- `finish`
  - Runs until the current function is finished
- `until`
  - When executed in a loop, continues the execution until the loop ends
- `info breakpoints`
  - Shows informations about all declared breakpoints

```
(gdb) info b
Num      Type           Disp Enb Address                  What
1        breakpoint      keep y   0x0000000000040059f in main at string-test.c:
8
breakpoint already hit 1 time
2        hw watchpoint  keep y                   i
breakpoint already hit 2 times
```

- `delete`
  - Deletes a breakpoint

# Valgrind

## Example

```
#include <stdlib.h>

void f(void)
{
    int *x = malloc(10 * sizeof(int));
    x[10] = 0;
}

int main(void)
{
    f();
    return 0;
}
```

# Valgrind

## Run

```
$ valgrind --leak-check=full ./valgrind_example
...
==31134== Invalid write of size 4
==31134==    at 0x108668: f (in /home/joel/work/ecs150/slides/tuto_gdb/code/valgrind_example)
==31134==    by 0x108679: main (in /home/joel/work/ecs150/slides/tuto_gdb/code/valgrind_example)
==31134== Address 0x51f0068 is 0 bytes after a block of size 40 alloc'd
==31134==    at 0x4C2CEDF: malloc (vg_replace_malloc.c:299)
==31134==    by 0x10865B: f (in /home/joel/work/ecs150/slides/tuto_gdb/code/valgrind_example)
==31134==    by 0x108679: main (in /home/joel/work/ecs150/slides/tuto_gdb/code/valgrind_example)
==31134==
==31134== HEAP SUMMARY:
==31134==    in use at exit: 40 bytes in 1 blocks
==31134==    total heap usage: 1 allocs, 0 frees, 40 bytes allocated
==31134==
==31134== 40 bytes in 1 blocks are definitely lost in loss record 1 of 1
==31134==    at 0x4C2CEDF: malloc (vg_replace_malloc.c:299)
==31134==    by 0x10865B: f (in /home/joel/work/ecs150/slides/tuto_gdb/code/valgrind_example)
==31134==    by 0x108679: main (in /home/joel/work/ecs150/slides/tuto_gdb/code/valgrind_example)
==31134==
==31134== LEAK SUMMARY:
==31134==    definitely lost: 40 bytes in 1 blocks
==31134==    indirectly lost: 0 bytes in 0 blocks
==31134==    possibly lost: 0 bytes in 0 blocks
==31134==    still reachable: 0 bytes in 0 blocks
==31134==    suppressed: 0 bytes in 0 blocks
==31134==
==31134== For counts of detected and suppressed errors, rerun with: -v
==31134== ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 0 from 0)
```

# C preprocessor: Macros

## Sequential values

### Example

```
#define ST_READY    0
#define ST_RUNNING  1
#define ST_BLOCKED  2

void show_state(struct uthread_tcb *tcb)
{
    printf("Thread %d: ", tcb->tid);
    switch(tcb->state) {
        case ST_READY:
            printf("Ready\n");
        case ST_RUNNING:
            printf("Running\n");
        case ST_BLOCKED:
            printf("Blocked\n");
    }
}
```

### After preprocessing

```
$ cpp file.c
...
```

```
...
void show_state(struct uthread_tcb *tcb)
{
    printf("Thread %d: ", tcb->tid);
    switch(tcb->state) {
        case 0:
            printf("Ready\n");
        case 1:
            printf("Running\n");
        case 2:
            printf("Blocked\n");
    }
}
```



# C preprocessor: Macros

## Versus enum

```
#define ST_READY    0
#define ST_RUNNING  1
#define ST_BLOCKED  2
```

```
enum state {
    ST_READY,
    ST_RUNNING,
    ST_BLOCKED,
};
```

## Pros

- Autonumbering
- If used in a switch-case, warn if switch-case isn't complete
- Debugger gets access to names

## Cons

- No control on the integer type the compiler chooses
  - (possible only in C++: `enum state : char { ... };`)

# C preprocessor: Macros

---

## Arbitrary values

### Example

```
#define HZ 100

#define TIMER_REG1 0x0
#define TIMER_REG2 0x4
#define TIMER_REG3 0x8

#define TIMER_REG1_MASK 0x8000FFFF

void configure_timer(int *timer_addr, int val)
{
    *(timer_addr + TIMER_REG1) = val & TIMER_REG1_MASK;
}
```

### Pros

- Recommended when names matter more than values
- And values are somewhat arbitrary (e.g., not in strict sequence)
  - Otherwise prefer `enum`

# C preprocessor: Macros

## Code replacement

### Example

```
#define twice(x) 2 * x

void f(int a)
{
    int b = twice(a);
}
```

### After preprocessing

```
void f(int a)
{
    int b = 2 * a;
}
```

### Pitfall #1

```
int c = twice(a + 1) * 3;
```

```
int c = 2 * a + 1 * 3;
```

- What's the value of `c`? Is it the value meant to be assigned?

### Solution

- Always surround your arguments with parenthesis!

```
#define twice(x) (2 * (x))
```

```
int c = (2 * (a + 1)) * 3;
```

# C preprocessor: Macros

## Pitfall #2

```
#define die_perror(x)  \
    perror(x);        \
    exit(1);

void f(void)
{
    int *a;

    a = malloc(10 * sizeof(int));
    if (a == NULL)
        die_perror("malloc");
    ...
}
```

- Any issues?

### After preprocessing

```
#define die_perror(x)  \
    perror(x);        \
    exit(1);

void f(void)
{
    int *a;

    a = malloc(10 * sizeof(int));
    if (a == NULL)
        perror("malloc");
        exit(1);
        ;
    ...
}
```

- Misleading, and logically incorrect

# C preprocessor: Macros

## Pitfall #2 (cont'd)

```
#define die_perror(x) \
{
    perror(x);
    exit(1);
}

void f(void)
{
    int *a;

    a = malloc(10 * sizeof(int));
    if (a == NULL)
        die_perror("malloc");
    ...
}
```

- Does this work better?

```
#define die_perror(x) \
{
    perror(x);
    exit(1);
}

void f(void)
{
    int *a;

    a = malloc(10 * sizeof(int));
    if (a == NULL)
        die_perror("malloc")
    ...
}
```

- Works but have to purposefully omit the semi-colon which is counter-intuitive

# C preprocessor: Macros

## Pitfall #2 (cont'd)

### Solution

```
#define die_perror(x) \
do { \
    perror(x); \
    exit(1); \
} while(0)

void f(void)
{
    int *a;

    a = malloc(10 * sizeof(int));
    if (a == NULL)
        die_perror("malloc");
    ...
}
```

- Surround your multiple statement code with `do { ... } while(0)`

# C preprocessor: Macros

## Pitfall #3

```
#define MAX(x,y) ((x) > (y) ? (x) : (y))
```

```
void f(int a, int b)
{
    int c = MAX(a, b);
}
```

```
int d = MAX(++a, b);
```

```
int d = ((++a) > (b) ? (++a) : (b))
```

- What's the value of d? Is it the value that was meant to be assigned?

## Solution

```
#define MAX(x,y) \
({ \
    typeof(x) _x = (x); \
    typeof(y) _y = (y); \
    (_x) > (_y) ? (_x) : (_y); \
})
```

- ({ to start a compound statement within an expression
- typeof(variable) replaced by type of variable

# C preprocessor: Macros

## Return value

```
void f(void)
{
    int b = ({int c = 2; printf("lol\n"); c;});
    printf("%d\n", b);
}
```

- Last statement of expression is expression's result

## Macro example

```
#define xmalloc(size) \
({ \
    void *x = malloc(size); \
    if (!x) { \
        perror("malloc"); \
        exit(1); \
    } \
    x; \
})

void f(void)
{
    int *a = xmalloc(10 * sizeof(int));
    ...
}
```

- Better to have `xmalloc` declared as a `static inline` (see next slide)



# C preprocessor: Macros

## Macro vs static inline

### static inline

```
static inline void* xmalloc(size_t size)
{
    void *x = malloc(size);
    if (!x) {
        perror("malloc");
        exit(1);
    }
    return x;
}
```

#### Pros

- Type checking
- Actual function

#### Cons

- Inlining is purely advisory (compiler is free to ignore it)

### Macro

```
#define WARN_IF(exp) \
do { \
    if (exp) \
        fprintf(stderr, \
            "Warning: " #exp "\n"); \
} while(0)

...
WARN_IF(my_ptr == NULL);
```

#### Pros

- No type checking (generic)
- Possibility to stringify arguments

#### Cons

- No type checking (or more difficult to implement)
- Difficult to debug
- Harder to edit when multi-line

# Multi-talented `printf()`

101

```
printf("Hello world!\n");
```

printf\_examples.c

```
$ ./a.out  
Hello world!
```

201

```
printf("%s from %c to Z, in %d minutes!\n", "printf", 'A', 45);
```

printf\_examples.c

```
$ ./a.out  
printf from A to Z, in 45 minutes!
```

`pow(101, 2)`

```
int i;  
printf("\b%n", &i);
```

printf\_examples.c

```
printf("%s\bD is \033[1;31m#%d\033[0m!\n", "UCB", i);
```

printf\_examples.c

```
$ ./a.out  
UCD is #1!
```

# printf(): an odyssey

---

## Fortran I

Special statement for building formatting descriptions:

```
WRITE OUTPUT TAPE 6, 601, IA, IB, IC, AREA
601 FORMAT (4H A= ,I5,5H B= ,I5,5H C= ,I5,8H AREA= ,F10.2,
+          13H SQUARE UNITS)
```

- (Approximate) translation in C:

```
printf(" A= %5d B= %5d C= %5d AREA= %10.2f SQUARE UNITS", a, b, c, area);
```

## BCPL

Printing and formatting are merged into a single statement:

```
WRITEF("%I2-QUEENS PROBLEM HAS %I5 SOLUTIONS*N", NUMQUEENS, COUNT)
```

- (Approximate) translation in C:

```
printf("%2d-queens problem has %5d solutions\n", numqueens, count);
```

# printf(): an odyssey

---

C

```
printf("Hello %s, you are %d years old\n", name, age);
```

## Trickle-down string formatting

Unix `printf`

```
$ printf "%s, stop lying; you're not %d!\n" Bob 21  
Bob, stop lying; you're not 21!
```

## Other languages...

awk, C++, Objective C, D, F#, G (LabVIEW), GNU MathProg, GNU Octave, Go, Haskell, J, Java (since version 1.5) and JVM languages (Clojure, Scala), Lua (string.format), Maple, MATLAB, Max (via the sprintf object), Mythryl, PARI/GP, Perl, PHP, Python (via % operator), R, Red/System, Ruby, Tcl (via format command), Transact-SQL (via xp\_sprintf), Vala.

# Why printf()?

## Output

```
$ cowsay "I love lectures about printf\!"
```

```
< I love lectures about printf! >
```

```
  \      ^__^
   (oo)\_____)
      (_____)  )\/\
        ||----w |
        ||     ||
```

```
$ cowthink -f vader "printf, I'm your father"
```

```
( printf, I'm your father )
```

```
  o      ^-^
   o    !oYo!
  o  /. /=\.\_____)\/\
      ##          ||
      ||----w |
      ||     ||
```

Cowth Vader

```
$ du /bin/* | sort -rn
```

```
20672 /bin/js52
19164 /bin/inkscape
19136 /bin/inkview
18720 /bin/node
18684 /bin/clementine
17452 /bin/mariabackup
17112 /bin/mysqld
16432 /bin/mysql_client_test_embedded
16332 /bin/mysql_embedded
16232 /bin/mysqltest_embedded
7600 /bin/gdb
...
```

## Introspection

```
Booting kernel from Legacy Image at 20080000 ...
```

```
Image Name:   Linux-2.6.37
Image Type:   ARM Linux Kernel Image (uncompressed)
Data Size:    1256880 Bytes = 1.2 MiB
Load Address: 20008000
Entry Point:  20008000
Verifying Checksum ... OK
Loading Kernel Image ... OK
```

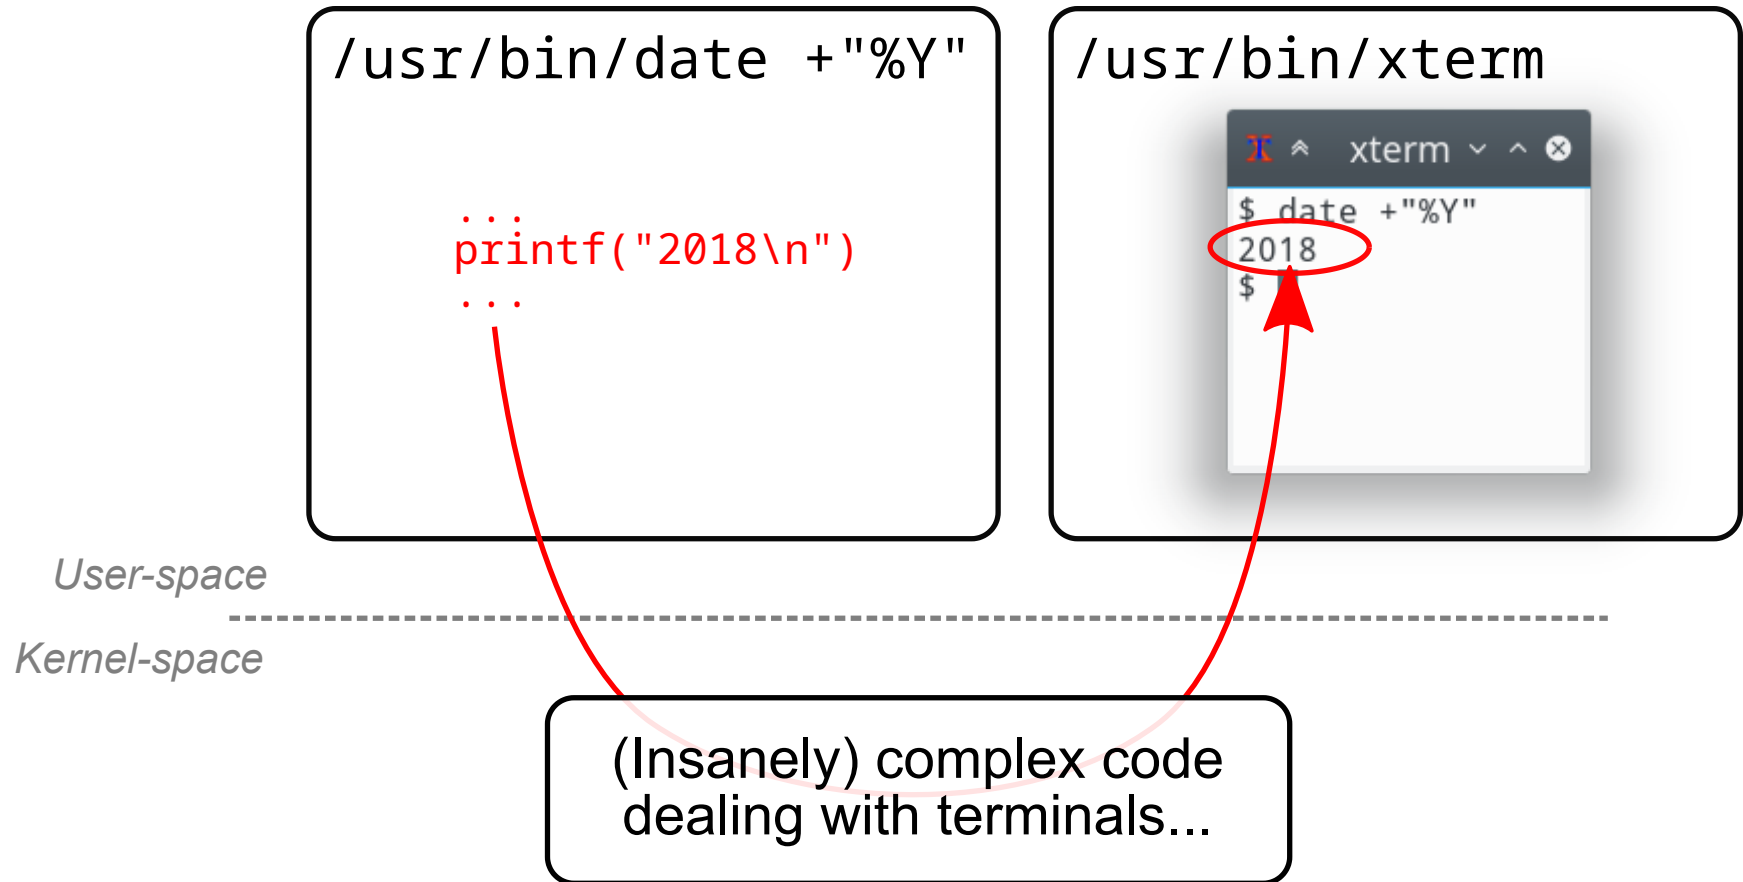
OK

Starting kernel ...

```
Uncompressing Linux... done, booting the kernel.
Linux version 2.6.37 (nkinar at matilda) (gcc version 4.3.5 (Buildroot 2011.02) ) #3 Sat Apr 2 17:28:21 CST 2011
CPU: ARM926EJ-S [41069265] revision 5 (ARMv5TEJ), cr=00053177
CPU: VIVT data cache, VIVT instruction cache
Machine: Atmel AT91SAM9RL-EK
Memory policy: ECC disabled, Data cache writeback
Clocks: CPU 200 MHz, master 100 MHz, main 12.000 MHz
Built 1 zonelists in Zone order, mobility grouping on.
Total pages: 16256
Kernel command line: console=ttyS0,115200
mtdparts=flash:10M(kernel),100M(root),-(storage) rw rootfstype=ubifs
PID hash table entries: 256 (order: -2, 1024 bytes)
Dentry cache hash table entries: 8192 (order: 3, 32768 bytes)
Inode-cache hash table entries: 4096 (order: 2, 16384 bytes)
Memory: 64MB = 64MB total
Memory: 62348k/62348k available, 3188k reserved, 0K highmem
...
```

# Tell me where you `printf()`!

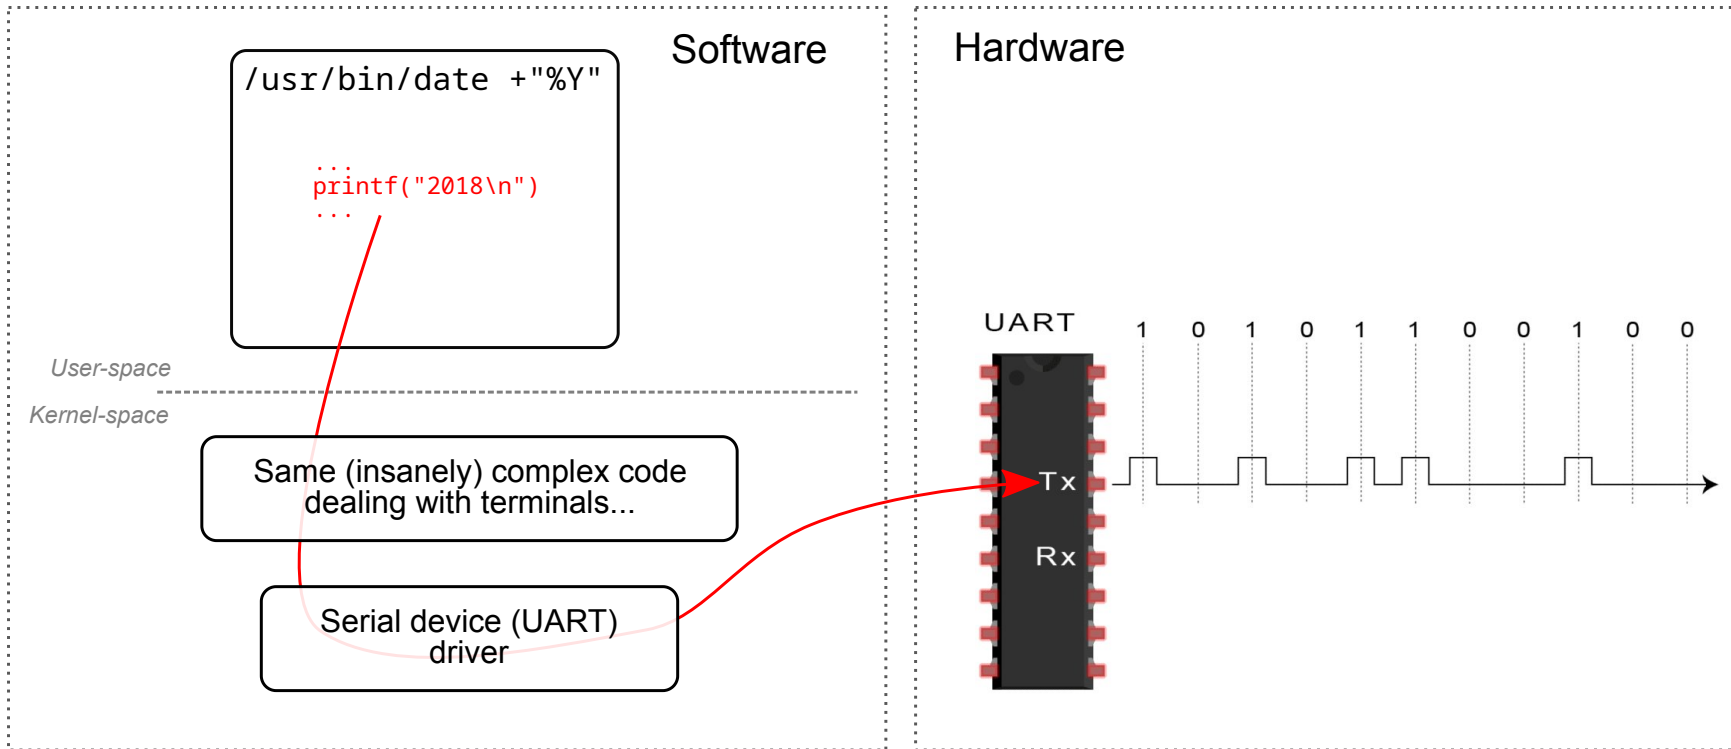
Typical GNU/Linux computer



*Ingo Molnár (Linux kernel core developer): "The tty layer is one of the very few pieces of kernel code that scares the hell out of me :-)"*

# Tell me where you `printf()`!

Embedded systems or when display is not available




- Characters are sent one by one over a serial port

# A serial port on every board?

---

My own phone!







# A serial port on every board?

## Playstation 4


PS4




Pro PS4




Slim PS4



## Linux on PS4



## PS4 controller



# How to print?

## putchar(): the cornerstone

```
/*
 * Code for IBM-PC x86: write character to COM1 serial port
 */
void putchar(char ch)
{
    /* Wait until the Transmitter Holding Register (THR) is empty. */
    while ((inb(COM1_PORT+COM_LSR) & LSR_THRE) == 0);

    /* Then output the character to the THR */
    outb(ch, COM1_PORT+COM_THR);
}
```

\* Copied from [NuttX](#) RTOS

# How to print?

No formatting is simply `puts ( )`

```
void my_printf(char *fmt)
{
    while (*fmt)
        putchar(*fmt++);
}

int main(void)
{
    /* printf 101 */
    my_printf("Hello world!\n");
}
```

printf\_puts.c

```
$ ./a.out
Hello world!
```

# How to print?

With formatting...



**Input:** `my_printf("%s from %c to Z, in %d minutes!\n", "printf", 'A', 45);`

**Output:** `printf from A to Z, in 45 minutes!\n`

- For each placeholder, need to retrieve next parameter
- Depending on placeholder:
  - `'%s'`: get string
  - `'%c'`: get character
  - `'%d'`: **convert** integer into characters

# Variadic functions

## Example: function prototype

```
#include <stdarg.h>    /* Macro/function definitions for variadic functions */
#include <stdio.h>

/*
 * Sum a variable number of integers
 * @count: the number of integer parameters
 *
 * Receive @count integers, sum them up and return the result
 */
int sum_ints(int count, ...)
{
    ...
}

int main(void)
{
    int res;

    /* sum up 3 integers: 10, 20 and 30 */
    printf("Sum is %d\n", sum_ints(3, 10, 20, 30));

    /* sum up 5 integers: 10, 20, 30, 40, 50 */
    printf("Sum is %d\n", sum_ints(5, 10, 20, 30, 40, 50));

    return 0;
}
```

variadic\_fcn.c

# Variadic functions

## Example: function implementation

```
#include <stdarg.h>    /* Macro/function definitions for variadic functions */
...
int sum_ints(int count, ...)
{
    va_list ap;
    int i, sum = 0;

    va_start(ap, count);          /* Init variable parameter list */

    for (i = 0; i < count; i++) {
        int n = va_arg(ap, int);  /* Get next integer parameter */
        sum += n;
    }

    va_end(ap);                  /* Clean up parameter list */

    return sum;
}
```

variadic\_fcn.c

# `printf()`: a simple implementation

---



Demo time: let's code!

# printf(): a simple implementation

---

## Before

```
#include <stdarg.h>
#include <stdio.h>

void my_printf(char *fmt, ...)
{
    /* ??? */
}

int main(void)
{
    /* printf 101 */
    my_printf("Hello world!\n");

    /* printf 201: basic placeholders */
    my_printf("%s from %c to Z, in %d minutes!\n",
               "printf", 'A', 0);

    return 0;
}
```

printf\_demo\_start.c



# printf(): a simple implementation

## After

```
void my_printf(char *fmt, ...)
{
    va_list ap;
    char c;
    int d;
    char *s;

    va_start(ap, fmt);

    while (*fmt) {
        /* Check if character is '%' */
        if (*fmt != '%') {
            /* If not, display without no processing and
             * continue to the next character
             */
            putchar(*fmt++);
            continue;
        }

        /* Skip '%' and get to the placeholder */
        fmt++;
        /* Distinguish different placeholders */
        switch(*fmt) {
            case 'c':
                /* Get character from arguments */
                c = va_arg(ap, int);
                putchar(c);
                break;
            case 's':
                /* Get string pointer from arguments */
                s = va_arg(ap, char*);
                /* Display each char from the string */
                while (*s) putchar(*s++);
                break;
            ...
        }
    }
}
```

printf\_demo\_end.c

```
        case 'd':
            /* Get integer from arguments */
            d = va_arg(ap, int);
            /* Translate integer into string
             * (via ASCII conversion) */
            char buf[10], *end = buf;
            do {
                *end = (d % 10) + '0';
                end++;
                d /= 10;
            } while (d);
            /* Display string */
            while (end != buf) {
                putchar(*--end);
            }
            break;
        }
        /* Skip placeholder and continue */
        fmt++;
    }

    va_end(ap);
}

int main(void)
{
    /* printf 101 */
    my_printf("Hello world!\n");

    /* printf 201: basic placeholders */
    my_printf("%s from %c to Z, in %d minutes!\n",
              "printf", 'A', 45);

    return 0;
}
```

printf\_demo\_end.c

```
printf( "Thanks!" );
```

---